

Phase 1 Foundation Encoder

Self-Supervised ECG Pre-Training - Evaluation Report

Masked ECG Reconstruction | 12-Lead | PTB-XL + Chapman + MIMIC-IV-ECG

Date:	30 September 2026	Test Records:	80 held-out samples
Checkpoint:	foundation_encoder_best.pth (~25.4 MB)	Device:	CUDA GPU (AMP Enabled)
Test Datasets:	PTB-XL (fold 10) + Chapman (10%) + MIMIC-IV-ECG (10%)	Mask Ratio:	60% Contiguous Block

Key Results at a Glance

Masked Patch MSE (Primary)

0.1278

Lower is better

Masked RMSE

0.3576

Lower is better

Mean Absolute Error (MAE)

0.2276

Lower is better

R-Squared (R2)

0.8715

Closer to 1.0 is better

Pearson Correlation (r)

0.9339

Closer to 1.0 is better

Signal-to-Noise Ratio (SNR)

8.91 dB

Higher is better

Peak SNR (PSNR)

14.95 dB

Higher is better

Overall Full-Sequence MSE

0.0968

Entire 10s sequence

NOTE: Pearson $r = 0.9339$ and $R^2 = 0.8715$ confirm strong linear agreement between reconstructed and original ECG signals on masked spans. The SNR of 8.91 dB indicates high-fidelity recovery despite 60% of the signal being erased as input.

1. Complete Metrics Summary (Masked Region)

Metric	Value	Description
Masked Patch MSE	0.127845	Primary reconstruction metric on 60% erased spans
Masked RMSE	0.357354	Square root of masked MSE, same units as signal amplitude
Mean Absolute Error	0.227601	Average absolute deviation across masked patches
R-Squared (R2)	0.871495	Proportion of waveform variance explained by model (1.0 = perfect)
Pearson Correlation (r)	0.933877	Linear correlation between reconstructed output and ground truth
Signal-to-Noise Ratio (SNR)	8.91 dB	Ratio of true signal power vs noise power on masked spans
Peak SNR (PSNR)	14.95 dB	Peak SNR; peak assumed = 2.0 (standard normalized range)
Overall Full-Seq. MSE	0.096765	MSE across entire 10-second sequence (masked + unmasked)
Unmasked Patch MSE	0.050206	Fidelity on the 40% visible context patches

INFO: All primary metrics except Overall MSE and Unmasked MSE are computed exclusively on the MASKED (erased) 60% region — the region the model must predict from context. This is the fair and standard evaluation protocol for masked autoencoders.

2. Model Architecture

Parameter	Value
Architecture	Transformer Foundation Encoder (ViT / MAE 1D Style)
Input Signal	12-lead ECG, 5000 samples (10 s @ 500 Hz)
Patch Size	50 samples (100 ms per token)
Number of Patches	100 patches per lead
Embedding Dimension	256 channels
Transformer Layers	8 Pre-LayerNorm Transformer Encoder blocks
Attention Heads	8 multi-head self-attention heads (32 dim/head)
Dropout (training)	0.1 for attention and feedforward layers
Positional Encoding	Learned 1D temporal positional embeddings
Total Parameter Count	~6.65 Million parameters (~25.4 MB checkpoint)

3. Pre-Training Configuration

Setting	Value
Objective	Masked ECG Reconstruction (MAE-style self-supervised learning)
Masking Strategy	Configurable Block Masking
Mask Ratio	60% of all temporal patches per lead
Optimizer	AdamW (beta=(0.9, 0.95))
Peak Learning Rate	2e-4
Weight Decay	0.05 (applied to weights, bias and LayerNorm exempt)
LR Schedule	Cosine Annealing (eta_min = 1e-6)
Batch Size	64 records per batch
Mixed Precision	Automatic Mixed Precision (AMP, FP16 forward / FP32 gradient)
Gradient Clipping	max_norm = 1.0
Unmasked Aux Loss Wt.	0.5 (prevents mean-collapse on masked spans)


4. Training Datasets

Dataset	Split Protocol	Description
PTB-XL	Train / Val / Test (fold 10 benchmark)	12-lead, 500 Hz, clinical diagnostic labels
Chapman	Train / Val / Test (10% hold-out)	12-lead, 500 Hz, arrhythmia annotations
MIMIC-IV-ECG	Train / Val / Test (10% patient-level)	12-lead, 500 Hz, large-scale ICU patient cohort

5. Per-Lead Masked Reconstruction MSE

Lead	Masked MSE	Clinical Performance Tier
I	0.091803	High Fidelity
II	0.098530	High Fidelity
III	0.123686	High Fidelity
aVR	0.079075	Best Performing Lead
aVL	0.124046	High Fidelity
aVF	0.114712	High Fidelity
V1	0.172106	Hardest Lead (High anatomical variance)
V2	0.132820	High Fidelity
V3	0.147028	High Fidelity
V4	0.146800	High Fidelity
V5	0.133232	High Fidelity
V6	0.170305	High Fidelity

Relative MSE Bar Chart (Masked Region):



NOTE: Lead aVR achieves the best reconstruction (MSE=0.0791). Lead V1 is the most challenging (MSE=0.1721). Limb leads (I, II, aVR) generally outperform precordial leads (V1-V6) due to their higher inter-lead correlation and simpler morphology.

6.1 ECG Reconstruction - Sample 1

Four panels shown for Lead II from a different held-out ECG record:

- * Panel 1: Original Ground Truth 10-second ECG
- * Panel 2: Masked Input (60% erased, red zones indicate masked patches)
- * Panel 3: Inpainted ECG (visible spans + model predictions filling the gaps)
- * Panel 4: Raw Model Output overlaid directly against Ground Truth



6.2 ECG Reconstruction - Sample 2

Four panels shown for Lead II from a different held-out ECG record:


- * Panel 1: Original Ground Truth 10-second ECG
- * Panel 2: Masked Input (60% erased, red zones indicate masked patches)
- * Panel 3: Inpainted ECG (visible spans + model predictions filling the gaps)
- * Panel 4: Raw Model Output overlaid directly against Ground Truth



6.3 ECG Reconstruction - Sample 3

Four panels shown for Lead II from a different held-out ECG record:


- * Panel 1: Original Ground Truth 10-second ECG
- * Panel 2: Masked Input (60% erased, red zones indicate masked patches)
- * Panel 3: Inpainted ECG (visible spans + model predictions filling the gaps)
- * Panel 4: Raw Model Output overlaid directly against Ground Truth



6.4 ECG Reconstruction - Sample 4

Four panels shown for Lead II from a different held-out ECG record:


- * Panel 1: Original Ground Truth 10-second ECG
- * Panel 2: Masked Input (60% erased, red zones indicate masked patches)
- * Panel 3: Inpainted ECG (visible spans + model predictions filling the gaps)
- * Panel 4: Raw Model Output overlaid directly against Ground Truth



6.5 ECG Reconstruction - Sample 5

Four panels shown for Lead II from a different held-out ECG record:

- * Panel 1: Original Ground Truth 10-second ECG
- * Panel 2: Masked Input (60% erased, red zones indicate masked patches)
- * Panel 3: Inpainted ECG (visible spans + model predictions filling the gaps)
- * Panel 4: Raw Model Output overlaid directly against Ground Truth



7. Key Observations

1. Strong Pearson Correlation ($r = 0.9339$)

A Pearson r above 0.93 indicates very strong linear agreement between the model's reconstructed waveforms and the ground truth on masked spans, confirming the model captures the correct morphology and amplitude scale.

2. High R-Squared ($R^2 = 0.8715$)

R^2 of 0.8715 means the model explains ~87.1% of the variance in masked ECG amplitudes purely from surrounding context — a strong result for a 60% masking ratio on 10-second signals.

3. Waveform Morphology Preserved

The reconstruction plots confirm P-waves, QRS complexes, and T-waves are correctly recovered in masked spans, demonstrating genuine temporal reasoning rather than trivial interpolation or mean-regression.

4. Limb Leads > Precordial Leads

Best lead: aVR (MSE=0.0791). Hardest lead: V1 (MSE=0.1721). Limb leads benefit from high inter-lead correlation; precordial leads capture more spatially localized cardiac activity that is harder to infer from context.

5. SNR = 8.91 dB / PSNR = 14.95 dB

An SNR of 8.91 dB and PSNR of 14.95 dB on masked regions indicates strong reconstruction fidelity for a pre-training stage. Fine-tuning on downstream tasks is expected to further improve representation quality.

6. No Mean-Collapse

The auxiliary unmasked loss (weight=0.5) prevents the model from predicting flat waveforms on masked regions. Evidence: unmasked MSE (0.050) is ~2.5x lower than masked MSE (0.128), confirming the difficulty gap is genuine.

8. Checkpoint Summary

File Name	Size	Purpose
foundation_encoder_best.pth	~25.4 MB	Best validation loss checkpoint (used for this evaluation)
foundation_encoder_phase1.pth	~25.4 MB	Final training epoch checkpoint

9. Next Steps - Phase 2 Fine-Tuning

This pre-trained foundation encoder serves as a frozen or fine-tuned backbone for downstream supervised tasks: Arrhythmia classification (AFib, bundle branch block), ST-elevation / ischaemia detection, and patient risk stratification.

TIP: Phase 1 self-supervised pre-training provides a strong, label-free initialisation. Phase 2 fine-tuning requires only a small fraction of labelled samples to achieve competitive downstream performance — the core benefit of foundation model pre-training.